

Practical Introduction to Kubernetes

Reorganised and adapted by Unai Arronategui for K3s,

from <https://kubernetesbyexample.com/>

Index

1. **Pods**
2. **Labels**
3. **Nodes**
4. **Logging**
5. **API server**
6. **Deployments**
7. **Services**
8. **Service discovery**
9. **Port forward**
10. **Namespaces**
11. **Volumes**
12. **Persistent volumes**
13. **Stateful sets**
14. **Jobs**
15. **Secrets**
16. **Init containers**

1. Pods

A pod is a collection of containers sharing a network and mount namespace and is the basic unit of deployment in Kubernetes. All containers in a pod are scheduled on the same node.

To launch a pod using the container [image](#) :

`quay.io/openshiftlabs/simpleservice:0.5.0`

and exposing a HTTP API on port 9876, execute:

```
$ kubectl run sise --image=quay.io/openshiftlabs/simpleservice:0.5.0 --port=9876
```

We can now see that the pod is running:

```
$ kubectl get pods -o wide
NAME                                READY    STATUS    RESTARTS    AGE
sise-574c4c79d6-rqxqg              1/1      Running   0           1m

$ kubectl top pod
NAME                                CPU(cores)    MEMORY(bytes)
sise-574c4c79d6-rqxqg              0m            12Mi

$ kubectl describe pod sise-574c4c79d6-rqxqg | grep IP:
IP:                                10.42.X.X
```

See section “Interaction with running Pods”, from “Kubectl Cheat Sheet”, to test from within the cluster, from other Pod

```
$ kubectl run -i --tty busybox1 --image=busybox -- sh
```

this sise pod is accessible via the pod IP `10.42.X.X`, which we’ve learned from the `kubectl describe` command above:

```
/ # wget -q -O - 10.42.X.X:9876/info
{"host": "10.42.X.X", "version": "0.5.0", "from": "10.42.Y.Y"}
```

If after login out from busybox1 , there is a need to log again in :

```
kubectl attach busybox1-7c8d57f869-6p5wq -c busybox1 -i -t
```

Note that basic `kubectl run` creates a [deployment](#) in Kubernetes version 1.17.4. To see them :

```
$ kubectl get deployment
```

So in order to get rid of the pod you have to execute `kubectl delete deployment` `busybox1` deployment needs to be deleted in the same way.

Using configuration file

You can also create a pod from a **configuration file**. In this case the [pod](#) is running the already known `simpleservice` image container from above, along with a generic CentOS container, so **2 containers** are started in **1 Pod** :

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/pods/pod.yaml

$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
twocontainers	2/2	Running	0	7s

Read the [configuration file](#) (in a web browser) to study how is composed.

Now we can exec into the CentOS container (named `shell`) and access the `simpleservice` container, in the same Pod, on **localhost**, as **they share network** configuration :

```
$ kubectl exec twocontainers -c shell -i -t -- bash
/# curl -s localhost:9876/info
{"host": "localhost:9876", "version": "0.5.0", "from": "127.0.0.1"}
```

Specify the `resources` field in the pod to influence how much CPU and/or RAM a container in a [pod](#) can use (here: 64MB of RAM and 0.5 CPUs):

```
$ kubectl create -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/pods/constraint-pod.yaml

$ kubectl describe pod constraintpod
...
Containers:
  size:
    ...
  Limits:
    cpu:      500m
    memory:   64Mi
  Requests:
    cpu:      500m
    memory:   64Mi
...
```

Read the [configuration file](#) (in a web browser) to see how resources have been modified, and learn more about resource constraints in Kubernetes via the docs [here](#) :

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-memory-resource/>

and [here](#):

<https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/>

To remove all the pods created, just run:

```
$ kubectl delete pod twocontainers  
$ kubectl delete pod constraintpod
```

To sum up, launching one or more containers (together) in Kubernetes is simple, however doing it directly as shown above comes with a serious limitation: you have to manually take care of keeping them running in case of a failure. A better way to supervise pods is to use [deployments](#), giving you much more control over the life cycle, including rolling out a new version.

2. Labels

Labels are the mechanism you use to organize Kubernetes objects. A label is a key-value pair with certain [restrictions](#) concerning length and allowed values but **without any pre-defined meaning**. So you're free to choose labels as you see fit, for example, to express environments such as 'this pod is running in production' or ownership, like 'department X owns that pod'.

Let's create a [pod](#) that initially has one label (env=development):

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/labels/pod.yaml

$ kubectl get pods --show-labels
NAME          READY   STATUS    RESTARTS   AGE   LABELS
labelex       1/1     Running   0          10m   env=development
```

In above get pods command note the --show-labels option that output the labels of an object in an additional column.

You can add a label to the pod as:

```
$ kubectl label pods labelex owner=michael

$ kubectl get pods --show-labels
NAME          READY   STATUS    RESTARTS   AGE   LABELS
labelex       1/1     Running   0          16m   env=development,owner=michael
```

To use a label for filtering, for example to list only pods that have an owner that equals michael, use the --selector option:

```
$ kubectl get pods --selector owner=michael
NAME          READY   STATUS    RESTARTS   AGE
labelex       1/1     Running   0          27m
```

The --selector option can be abbreviated to -l, so to select pods that are labelled with env=development, do:

```
$ kubectl get pods -l env=development
NAME          READY   STATUS    RESTARTS   AGE
labelex       1/1     Running   0          27m
```

Oftentimes, Kubernetes objects also support set-based selectors. Let's launch [another pod](#) that has two labels (env=production and owner=michael):

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/
```

Now, let's list all pods that are either labelled with `env=development` or with `env=production`:

```
$ kubectl get pods -l 'env in (production, development)'
NAME           READY   STATUS    RESTARTS   AGE
labelex         1/1     Running   0           43m
labelexother    1/1     Running   0           3m
```

Other verbs also support label selection, for example, you could remove both of these pods with:

```
$ kubectl delete pods -l 'env in (production, development)'
```

Beware that this will destroy any pods with those labels.

You can also delete them directly, via their names, with:

```
$ kubectl delete pods labelex
$ kubectl delete pods labelexother
```

Note that **labels** are **not restricted** to **pods**. In fact you can apply them to all sorts of objects, such as nodes or services.

3. Nodes

In Kubernetes, nodes are the (virtual) machines where your workloads in shape of pods run. As a developer you typically don't deal with nodes directly, however as an admin you might want to familiarize yourself with node [operations](#).

To list available nodes in your cluster (note that the output will depend on the environment you're using. This example is using the [OpenShift Playground](#)):

```
$ kubectl get nodes
NAME      STATUS    ROLES    AGE     VERSION
m         Ready     master   42m     v1.17.4+k3s1
w2        Ready     <none>   41m     v1.17.4+k3s1
w1        Ready     <none>   41m     v1.17.4+k3s1
```

```
$ kubectl top nodes
NAME      CPU(cores)   CPU%    MEMORY(bytes)  MEMORY%
m         130m         13%     488Mi          50%
w1        48m          4%      339Mi          35%
w2        29m          2%      267Mi          27%
```

One interesting task, from a developer point of view, is to make Kubernetes schedule a pod on a certain node. For this, we first need to label the node we want to target:

```
$ kubectl label nodes crc-rk2fc-master-0 shouldrun=here
node/crc-rk2fc-master-0 labeled
```

Now we can create a [pod](#) that gets scheduled on the node with the label shouldrun=here:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/nodes/pod.yaml

$ kubectl get pods -o wide
NAME                READY  STATUS   RESTARTS  AGE    IP           NODE    NOMINATED NODE   READINESS GATES
onspecificnode      1/1    Running   0          2m31s  10.42.1.5    w1      <none>            <none>
```

To learn more about a specific node, crc-rk2fc-master-0 in our case, do:

```
$ kubectl describe node m
Name:               m
Roles:              master
Labels:             beta.kubernetes.io/arch=amd64
                   beta.kubernetes.io/instance-type=k3s
                   beta.kubernetes.io/os=linux
                   k3s.io/hostname=m
                   k3s.io/internal-ip=192.168.1.90
```

```

kubernetes.io/arch=amd64
kubernetes.io/hostname=m
kubernetes.io/os=linux
node-role.kubernetes.io/master=true
node.kubernetes.io/instance-type=k3s
Annotations:  alpha.kubernetes.io/provided-node-ip: 192.168.1.90
              flannel.alpha.coreos.com/backend-data:
{"VtepMAC":"fe:2d:45:28:d4:53"}
              flannel.alpha.coreos.com/backend-type: vxlan
              flannel.alpha.coreos.com/kube-subnet-manager: true
              flannel.alpha.coreos.com/public-ip: 192.168.1.90
              k3s.io/node-args:
                ["server","--token","*****","--flannel-iface","enp0s8","--bind-
address","192.168.1.90","--node-ip","192.168.1.90","--node-name","m","--...
              k3s.io/node-config-hash:
NVVADHZNVLE3W20QXJNHFTDR5Y46UG543H2H6XGKWEVZVIPGWDJQ====
              k3s.io/node-env:
{"K3S_DATA_DIR":"/var/lib/rancher/k3s/data/6a3098e6644f5f0dbfe14e5efa99bb8fdf60d63cae89f
dfffd71b7de11a1f1430"}
              node.alpha.kubernetes.io/ttl: 0
              volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Fri, 24 Apr 2020 12:43:00 +0200
Taints:          k3s-controlplane=true:NoExecute
Unschedulable:   false
Lease:
  HolderIdentity:  m
  AcquireTime:     <unset>
  RenewTime:       Fri, 24 Apr 2020 13:29:55 +0200
Conditions:
  Type            Status  LastHeartbeatTime               LastTransitionTime
Reason           Message
  ----           -
NetworkUnavailable False   Fri, 24 Apr 2020 12:43:15 +0200   Fri, 24 Apr 2020
12:43:15 +0200   FlannelIsUp                     Flannel is running on this node
MemoryPressure   False   Fri, 24 Apr 2020 13:28:13 +0200   Fri, 24 Apr 2020
12:43:00 +0200   KubeletHasSufficientMemory      kubelet has sufficient memory available
DiskPressure     False   Fri, 24 Apr 2020 13:28:13 +0200   Fri, 24 Apr 2020
12:43:00 +0200   KubeletHasNoDiskPressure        kubelet has no disk pressure
PIDPressure      False   Fri, 24 Apr 2020 13:28:13 +0200   Fri, 24 Apr 2020
12:43:00 +0200   KubeletHasSufficientPID         kubelet has sufficient PID available
Ready            True    Fri, 24 Apr 2020 13:28:13 +0200   Fri, 24 Apr 2020
12:43:10 +0200   KubeletReady                    kubelet is posting ready status. AppArmor
enabled
Addresses:
  InternalIP: 192.168.1.90
  Hostname:   m

```



```

Capacity:
  cpu:          1
  ephemeral-storage: 10098432Ki
  hugepages-2Mi: 0
  memory:       984228Ki
  pods:         110
Allocatable:
  cpu:          1
  ephemeral-storage: 9823754642
  hugepages-2Mi: 0
  memory:       984228Ki
  pods:         110
System Info:
  Machine ID:      caad551e95584e119f46ac9e133d7187
  System UUID:     1A2BF7FC-C6D9-45D3-803C-8DF306BC8446
  Boot ID:         38c5d79b-a990-4c74-8d8b-1487f3cff459
  Kernel Version:  4.15.0-96-generic
  OS Image:        Ubuntu 18.04.4 LTS
  Operating System: linux
  Architecture:    amd64
  Container Runtime Version: containerd://1.3.3-k3s2
  Kubelet Version: v1.17.4+k3s1
  Kube-Proxy Version: v1.17.4+k3s1
PodCIDR:          10.42.0.0/24
PodCIDRs:         10.42.0.0/24
ProviderID:       k3s://m
Non-terminated Pods: (0 in total)
  Namespace          Name          CPU Requests  CPU Limits  Memory Requests  Memory
Limits  AGE
  -----
  -----
Allocated resources:
  (Total limits may be over 100 percent, i.e., overcommitted.)
  Resource           Requests  Limits
  -----
  cpu                0 (0%)   0 (0%)
  memory             0 (0%)   0 (0%)
  ephemeral-storage  0 (0%)   0 (0%)
Events:
  Type    Reason          Age          From          Message
  ----    -
  Normal  Starting        47m          kubelet, m     Starting kubelet.
  Warning InvalidDiskCapacity 47m          kubelet, m     invalid capacity 0
on image filesystem
  Normal  NodeHasSufficientMemory 47m (x2 over 47m) kubelet, m     Node m status is
now: NodeHasSufficientMemory
  Normal  NodeHasNoDiskPressure 47m (x2 over 47m) kubelet, m     Node m status is

```

```

now: NodeHasNoDiskPressure
  Normal    NodeHasSufficientPID    47m (x2 over 47m)    kubelet, m    Node m status is
now: NodeHasSufficientPID
  Normal    NodeAllocatableEnforced    47m                kubelet, m    Updated Node
Allocatable limit across pods
  Normal    Starting                    47m                kube-proxy, m    Starting kube-
proxy.
  Normal    NodeReady                   46m                kubelet, m    Node m status is
now: NodeReady

```

Note that there are more sophisticated methods than shown above, such as using affinity, to [assign pods to nodes](#) :

<https://kubernetes.io/docs/concepts/configuration/assign-pod-node/>

and depending on your use case, you might want to check those out as well.

4. Logging

Logging is one option to understand what is going on inside your applications and the cluster at large. Basic logging in Kubernetes makes the output a container produces available, which is a good use case for debugging. More advanced [setups](#) consider logs across nodes and store them in a central place, either within the cluster or via a dedicated (cloud-based) service.

Let's create a [pod](#) called `logme` that runs a container writing to `stdout` and `stderr`:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/logging/pod.yaml
```

To view the five most recent log lines of the `gen` container in the `logme` pod, execute:

```
$ kubectl logs --tail=5 logme -c gen
Thu Apr 27 11:34:40 UTC 2017
Thu Apr 27 11:34:41 UTC 2017
Thu Apr 27 11:34:41 UTC 2017
Thu Apr 27 11:34:42 UTC 2017
Thu Apr 27 11:34:42 UTC 2017
```

To stream the log of the `gen` container in the `logme` pod (like `tail -f`), do:

```
$ kubectl logs -f --since=10s logme -c gen
Thu Apr 27 11:43:11 UTC 2017
Thu Apr 27 11:43:11 UTC 2017
Thu Apr 27 11:43:12 UTC 2017
Thu Apr 27 11:43:12 UTC 2017
Thu Apr 27 11:43:13 UTC 2017
...
```

Note that if you wouldn't have specified `--since=10s` in the above command, you would have gotten all log lines from the start of the container.

You can also view logs of pods that have already completed their lifecycle. For this we create a [pod](#) called `oneshot` that counts down from 9 to 1 and then exits. Using the `-p` option you can print the logs for previous instances of the container in a pod:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/logging/oneshotpod.yaml
$ kubectl logs -p oneshot -c gen
9
8
7
6
```

5
4
3
2
1

You can remove the created pods with:

```
$ kubectl delete pod/logme pod/oneshot
```

5. API Server access

Sometimes it's useful or necessary to directly [access the Kubernetes API server](#), for exploratory or testing purposes.

In order to do this, one option is to proxy the API to your local environment (your host operating system), using:

```
$ kubectl proxy --port=8080
Starting to serve on 127.0.0.1:8080
```

Now you can query the API (in a separate terminal session) like so:

```
$ curl http://localhost:8080/api/v1
{
  "kind": "APIResourceList",
  "groupVersion": "v1",
  "resources": [
    {
    ...
    {
      "name": "services/status",
      "singularName": "",
      "namespaced": true,
      "kind": "Service",
      "verbs": [
        "get",
        "patch",
        "update"
      ]
    }
  ]
}
```

Alternatively, without proxying, you can use `kubectl` directly as follows to achieve the same:

```
$ kubectl get --raw /api/v1
```

Further, if you want to explore the supported API versions and/or resources, you can use the following commands:

```
$ kubectl api-versions
admissionregistration.k8s.io/v1beta1
...
```

v1

```
$ kubectl api-resources
```

NAME		SHORTNAMES	APIGROUP	NAMESPACED
KIND				
bindings	true	Binding		
componentstatuses		cs	false	ComponentStatus
configmaps		cm	true	ConfigMap
...				

```
hljs.initHighlightingOnLoad(); renderMathInElement(document.body); var doNotTrack =
false;
if (!doNotTrack) {
    (function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
    (i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new Date();a=s.createElement(o),
    m=s.getElementsByTagName(o)
[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
    })(window,document,'script','https://www.google-analytics.com/analytics.js','ga');
    ga('create', 'UA-34425776-2', 'auto');

    ga('send', 'pageview');
} hljs.initHighlightingOnLoad(); renderMathInElement(document.body); var doNotTrack =
false;
if (!doNotTrack) {
    (function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
    (i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new Date();a=s.createElement(o),
    m=s.getElementsByTagName(o)
[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
    })(window,document,'script','https://www.google-analytics.com/analytics.js','ga');
    ga('create', 'UA-34425776-2', 'auto');

    ga('send', 'pageview');
} hljs.initHighlightingOnLoad(); renderMathInElement(document.body); var doNotTrack =
false;
if (!doNotTrack) {
    (function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){
    (i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new Date();a=s.createElement(o),
    m=s.getElementsByTagName(o)
[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)
    })(window,document,'script','https://www.google-analytics.com/analytics.js','ga');
    ga('create', 'UA-34425776-2', 'auto');
```

```
ga('send', 'pageview');  
}
```

6. Deployments

A **deployment** is a **supervisor** for **pods**, giving you fine-grained control over how and when a new pod version is rolled out as well as rolled back to a previous state.

Let's create a **deployment** called **sise-deploy** that supervises two replicas of a pod as well as a replica set:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/deployments/d09.yaml
```

You can have a look at the deployment, as well as the the replica set and the pods the deployment looks after like so:

```
$ kubectl get deploy
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
sise-deploy   2         2         2            2           10s

$ kubectl get rs
NAME                DESIRED   CURRENT   READY   AGE
sise-deploy-3513442901  2         2         2       19s

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
sise-deploy-3513442901-cndsx       1/1     Running   0          25s
sise-deploy-3513442901-sn74v       1/1     Running   0          25s
```

Note the **naming** of the **pods** and **replica set**, **derived** from the deployment name.

At this point in time the sise containers running in the pods are configured to return the version 0.9. Let's verify that from within the cluster (using *kubectl describe* first to get the IP of one of the pods and, next, using a *busybox* pod as in the beginning):

```
/ # wget -q -O - 10.42.X.X:9876/info
{"host": "10.42.X.X:9876", "version": "0.9", "from": "10.42.Y.Y"}
```

Let's now see what happens if we change that version to 1.0 in an updated **deployment**:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/deployments/d10.yaml
deployment "sise-deploy" configured
```

Note that you could have used `kubectl edit deploy/sise-deploy` alternatively to achieve the same by manually editing the deployment.

What we now see is the rollout of two new pods with the updated version 1.0 as well as the two old pods with version 0.9 being terminated:

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
sise-deploy-2958877261-nfv28	1/1	Running	0	25s
sise-deploy-2958877261-w024b	1/1	Running	0	25s
sise-deploy-3513442901-cndsx	1/1	Terminating	0	16m
sise-deploy-3513442901-sn74v	1/1	Terminating	0	16m

Also, a new replica set has been created by the deployment:

```
$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
sise-deploy-2958877261	2	2	2	4s
sise-deploy-3513442901	0	0	0	24m

Note that during the deployment you can check the progress using `kubectl rollout status deploy/sise-deploy`.

To verify that if the new 1.0 version is really available, you need to be **aware** that **pods** has **changed** the **name** and **IP**, so **again using** `kubectl describe` get the IP of one of the new pods, and access to them with `busybox` Pod:

```
/ # wget -q -O - 10.42.X.X:9876/info  
{"host": "10.42.X.X:9876", "version": "1.0", "from": "10.42.Y.Y"}
```

A history of all deployments is available via:

```
$ kubectl rollout history deploy/sise-deploy  
deployments "sise-deploy"  
REVISION      CHANGE-CAUSE  
1              <none>  
2              <none>
```

If there are problems in the deployment Kubernetes will automatically roll back to the previous version, however you can also explicitly roll back to a specific revision, as in our case to revision 1 (the original pod version):

```
$ kubectl rollout undo deploy/sise-deploy --to-revision=1  
deployment "sise-deploy" rolled back  
  
$ kubectl rollout history deploy/sise-deploy  
deployments "sise-deploy"  
REVISION      CHANGE-CAUSE  
2              <none>  
3              <none>  
  
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
sise-deploy-3513442901-ng8fz	1/1	Running	0	1m
sise-deploy-3513442901-s8q4s	1/1	Running	0	1m

At this point in time we're back at where we started, with two new pods serving again version 0.9.

Finally, to clean up, we remove the deployment and with it the replica sets and pods it supervises:

```
$ kubectl delete deploy sise-deploy
deployment "sise-deploy" deleted
```

See also the [docs](#):

<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>

for more options on deployments and when they are triggered.

7. Services

A service is an abstraction for pods, providing a stable, so called virtual IP (VIP) address. While pods may come and go and with it their IP addresses, a service **allows clients to reliably connect** to the containers running in the pod using the VIP. The virtual in VIP means it is not an actual IP address connected to a network interface, but its purpose is purely to forward traffic to one or more pods. Keeping the mapping between the VIP and the pods up-to-date is the job of [kube-proxy](#), a process that runs on every node, which queries the API server to learn about new services in the cluster.

Let's create a pod supervised by an [ReplicationController \(RC\)](#) and a [service](#) along with it:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/services/rc.yaml

$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/services/svc.yaml
```

Now we have the supervised pod running:

```
$ kubectl get pods -l app=sise
NAME          READY   STATUS    RESTARTS   AGE
rcsise-6nq3k  1/1     Running   0           57s

$ kubectl describe pod rcsise-6nq3k | less
Name:          rcsise-6nq3k
Namespace:     default
Security Policy: restricted
Node:          localhost/192.168.99.100
Start Time:    Tue, 25 Apr 2017 14:47:45 +0100
Labels:        app=sise
Status:        Running
IP:            10.42.2.11
Controllers:   ReplicationController/rcsise
Containers:
...
```

You can, from within the cluster, access the pod directly via its assigned IP 10.42.2.11:

```
/ # wget -q -O - 10.42.2.11:9876/info
{"host": "10.42.2.11:9876", "version": "0.5.0", "from": "10.42.Y.Y"}
```

This is however, as mentioned above, **not advisable** since the **IPs** assigned to **pods** **may change**. **Hence**, enter the **simpleservice** we've created:

```
$ kubectl get svc
NAME                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
simpleservice        172.30.228.255   <none>           80/TCP           5m

$ kubectl describe svc simpleservice | less
Name:                simpleservice
Namespace:           default
Labels:              <none>
Selector:            app=sise
Type:                ClusterIP
IP:                  10.43.170.9
Port:                <unset> 80/TCP
Endpoints:           10.42.3.2:9876
Session Affinity:    None
No events.
```

The **service** keeps track of the **pods** it forwards traffic to **through** the **label**, in our case `app=sise`.

To observe their behavior from within the cluster, we start, again, the *busybox* Pod. And from there we can now access **simpleservice** with the **IP** of the **service**, not the Pod IP:

```
/ # wget -q -O - 10.43.48.66:80/info
{"host": "10.43.48.66:80", "version": "0.5.0", "from": "10.42.3.8"}
```

What makes the Virtual IP (VIP) 10.43.48.66 forward the traffic to the pod? The answer is: [IPtables](#), which is essentially a long list of rules that tells the Linux kernel what to do with a certain IP package.

Looking at the rules that concern our service (executed on a cluster node - *vagrant ssh w1* -) yields:

```
vagrant@w1:~$ sudo iptables-save | egrep "simpleservice|KUBE-SVC-EZC6WLOVQADP4IAW|
KUBE-:KUBE-SEP-BKRJGF27PSUUPB5L - [0:0]
:KUBE-SVC-EZC6WLOVQADP4IAW - [0:0]
-A KUBE-SEP-BKRJGF27PSUUPB5L -s 10.42.3.2/32 -j KUBE-MARK-MASQ
-A KUBE-SEP-BKRJGF27PSUUPB5L -p tcp -m tcp -j DNAT --to-destination 10.42.3.2:9876
-A KUBE-SERVICES ! -s 10.42.0.0/16 -d 10.43.170.9/32 -p tcp -m comment --comment
"default/simpleservice: cluster IP" -m tcp --dport 80 -j KUBE-MARK-MASQ
-A KUBE-SERVICES -d 10.43.170.9/32 -p tcp -m comment --comment "default/simpleservice:
cluster IP" -m tcp --dport 80 -j KUBE-SVC-EZC6WLOVQADP4IAW
-A KUBE-SVC-EZC6WLOVQADP4IAW -j KUBE-SEP-BKRJGF27PSUUPB5L
```

Above you can see the four rules that kube-proxy has thankfully added to the routing table, essentially stating that TCP traffic to 10.43.170.9:80 should be forwarded to

10.42.3.2:9876, which is our pod.

Let's now add a second pod by scaling up the RC supervising it:

```
$ kubectl scale --replicas=2 rc/rcsise
replicationcontroller "rcsise" scaled
```

```
$ kubectl get pods -l app=sise
```

NAME	READY	STATUS	RESTARTS	AGE
rcsise-6nq3k	1/1	Running	0	15m
rcsise-nv8zm	1/1	Running	0	5s

When we now check the relevant parts of the routing table again, now in machine [vagrant@m](#), we notice the addition of a bunch of IPtables rules:

```
vagrant@m:~$ sudo iptables-save | egrep "simpleservice|KUBE-SVC-EZC6WLOVQADP4IAW|KUBE-
UBE-SEP-VU2JG0HHBAZT2PP0"
:KUBE-SEP-BKRJGF27PSUUPB5L - [0:0]
:KUBE-SEP-VU2JG0HHBAZT2PP0 - [0:0]
:KUBE-SVC-EZC6WLOVQADP4IAW - [0:0]
-A KUBE-SEP-BKRJGF27PSUUPB5L -s 10.42.3.2/32 -j KUBE-MARK-MASQ
-A KUBE-SEP-BKRJGF27PSUUPB5L -p tcp -m tcp -j DNAT --to-destination 10.42.3.2:9876
-A KUBE-SEP-VU2JG0HHBAZT2PP0 -s 10.42.2.3/32 -j KUBE-MARK-MASQ
-A KUBE-SEP-VU2JG0HHBAZT2PP0 -p tcp -m tcp -j DNAT --to-destination 10.42.2.3:9876
-A KUBE-SERVICES ! -s 10.42.0.0/16 -d 10.43.170.9/32 -p tcp -m comment --comment
"default/simpleservice: cluster IP" -m tcp --dport 80 -j KUBE-MARK-MASQ
-A KUBE-SERVICES -d 10.43.170.9/32 -p tcp -m comment --comment "default/simpleservice:
cluster IP" -m tcp --dport 80 -j KUBE-SVC-EZC6WLOVQADP4IAW
-A KUBE-SVC-EZC6WLOVQADP4IAW -m statistic --mode random --probability 0.500000000000 -j
KUBE-SEP-VU2JG0HHBAZT2PP0
-A KUBE-SVC-EZC6WLOVQADP4IAW -j KUBE-SEP-BKRJGF27PSUUPB5L
```

In above routing table listing we see rules for the newly created pod serving at 10.42.2.3:9876 as well as an additional rule:

```
-A KUBE-SVC-EZC6WLOVQADP4IAW -m statistic --mode random --probability 0.500000000000 -j
KUBE-SEP-VU2JG0HHBAZT2PP0
```

This causes the traffic to the service being equally split between our two pods by invoking the statistics module of IPtables.

You can remove all the resources created by doing:

```
$ kubectl delete svc simpleservice
```

```
$ kubectl delete rc rcsise
```

8. Service Discovery

Service discovery is the process of figuring out how to connect to a [service](#). While there is a service discovery option based on [environment variables](#) available, the DNS-based service discovery is preferable. Note that DNS is a [cluster add-on](#) so make sure your Kubernetes distribution provides for one or install it yourself.

Let's create a [service](#) named `thesvc` and an [RC](#) supervising some pods along with it:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/sd/rc.yaml

$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/sd/svc.yaml
```

Now we want to connect to the `thesvc` service from within the cluster, say, from another service. To simulate this, we create a [jump pod](#) in the same namespace (default, since we didn't specify anything else):

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/sd/jumpod.yaml
```

The DNS add-on will make sure that our service `thesvc` is available via the Fully Qualified Domain Name (DNS) `thesvc.default.svc.cluster.local` from other pods in the cluster.

We can directly connect to and consume the service (in the same namespace) like so:

```
$ kubectl exec -it jumpod -c shell -- curl http://thesvc/info
{"host": "thesvc", "version": "0.5.0", "from": "10.42.1.5"}
```

Note that the IP address `172.17.0.5` above is the cluster-internal IP address of the jump pod.

To access a service that is deployed in a different namespace than the one you're accessing it from, use a FQDN in the form `$SVC.$NAMESPACE.svc.cluster.local`.

Let's see how that works by creating:

1. a [namespace](#) other
2. a [service](#) `thesvc` in namespace other
3. an [RC](#) supervising the pods, also in namespace other

If you're not familiar with namespaces, check out the [namespace examples](#) first.

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/sd/other-ns.yaml

$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/sd/other-rc.yaml

$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/sd/other-svc.yaml
```

We're now in the position to consume the service thesvc in namespace other from the default namespace (again via the jump pod):

```
$ kubectl exec -it jumpod -c shell -- curl http://thesvc.other/info
{"host": "thesvc.other", "version": "0.5.0", "from": "10.42.1.5"}
```

Summing up, DNS-based service discovery provides a flexible and generic way to connect to services across the cluster.

You can destroy all the resources created with:

```
$ kubectl delete pods jumpod

$ kubectl delete svc thesvc

$ kubectl delete rc rcsise

$ kubectl delete ns other
```

Keep in mind that **removing a namespace will destroy every resource inside.**

9. Port Forward

Kubernetes port forward by example

In the context of developing apps on Kubernetes it is often useful to quickly access a service from your local environment without exposing it using, for example, a load balancer or an ingress resource. In this case you can use [port forwarding](#).

Let's create an [app](#) consisting of a deployment and a service called `simpleservice`, serving on port 80:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/pf/app.yaml
```

Let's say we want to access the `simpleservice` service from the local environment, say, your laptop, on port 8080. So we forward the traffic as follows:

```
$ kubectl port-forward service/simpleservice 8080:80
Forwarding from 127.0.0.1:8080 -> 9876
Forwarding from [::1]:8080 -> 9876
```

We can see from above that the traffic gets eventually routed through the service to the pod serving on port 9876.

Now we can invoke the service locally like so (using a separate terminal session):

```
$ curl localhost:8080/info
{"host": "localhost:8080", "version": "0.5.0", "from": "127.0.0.1"}
```

Remember that port forwarding is not meant for production traffic but for development and experimentation.

You can remove the `simpleservice` with

```
$ kubectl delete service/simpleservice
$ kubectl delete deployment sise-deploy
```


10. Namespaces

Namespaces provide a scope for Kubernetes resources, carving up your cluster in smaller units. You can think of it as a workspace you're sharing with other users. Many resources such as pods and services are namespaced, while some, for example, nodes are not namespaced (but cluster-wide). As a developer you'd usually use an assigned namespace, however admins may wish to manage them, for example to set up access control or resource quotas.

Let's list all namespaces (note that the output will depend on the environment you're using):

```
$ kubectl get ns
NAME                STATUS    AGE
default             Active    20m
kube-system         Active    20m
kube-public         Active    20m
kube-node-lease     Active    20m
...
```

You can learn more about a namespace using the describe verb, for example:

```
$ kubectl describe ns default
Name:          default
Labels:        <none>
Annotations:   <none>
Status:        Active

No resource quota.

No LimitRange resource.
```

Let's now create a new [namespace](#) called test now:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/ns/ns.yaml
namespace "test" created

$ kubectl get ns
NAME                STATUS    AGE
default             Active    24m
kube-system         Active    24m
kube-public         Active    24m
kube-node-lease     Active    24m
test                Active    29s
```

Alternatively, we could have created the namespace using the `kubectl create namespace test` command.

To launch a [pod](#) in the newly created namespace test, do:

```
$ kubectl apply --namespace=test -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/ns/pod.yaml
```

Note that using above method the namespace becomes a runtime property, that is, you can deploy the same pod or service, etc. into multiple namespaces (for example: dev and prod). Hard-coding the namespace directly in the metadata section like shown in the following is possible but causes less flexibility when deploying your apps:

```
apiVersion: v1
kind: Pod
metadata:
  name: podintest
  namespace: test
```

To list namespaced objects such as our pod podintest, run the following command:

```
$ kubectl get pods --namespace=test
```

NAME	READY	STATUS	RESTARTS	AGE
podintest	1/1	Running	0	16s

You can remove the namespace (**and everything inside**) with:

```
$ kubectl delete ns test
```

If you're an admin, you might want to check out the [docs](#) :

<https://kubernetes.io/docs/tasks/administer-cluster/namespaces/>

for more info how to handle namespaces.

11. Volumes

A Kubernetes volume is essentially a directory accessible to all containers running in a pod. In contrast to the container-local filesystem, the data in volumes is preserved across container restarts. The medium backing a volume and its contents are determined by the volume type:

- node-local types such as *emptyDir* or *hostPath*
- file-sharing types such as *nfs*
- cloud provider-specific types like *awsElasticBlockStore*, *azureDisk*, or *gcePersistentDisk*
- distributed file system types, for example *glusterfs* or *cephfs*
- special-purpose types like *secret*, *gitRepo*

A special type of volume is *PersistentVolume*, which we will cover elsewhere.

Let's create a [pod](#) with two containers that use an *emptyDir* volume to exchange data:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/volumes/pod.yaml

$ kubectl describe pod sharevol
Name:                sharevol
Namespace:           default
...
Volumes:
  xchange:
    Type:             EmptyDir (a temporary directory that shares a pod's lifetime)
    Medium:
```

We first exec into one of the containers in the pod, *c1*, check the volume mount and generate some data:

```
$ kubectl exec -it sharevol -c c1 -- bash
[root@sharevol /]# mount | grep xchange
/dev/vda3 on /tmp/xchange type xfs (rw,relatime,seclabel,attr2,inode64,rjquota)
[root@sharevol /]# echo 'some data' > /tmp/xchange/data
```

When we now exec into c2, the second container running in the pod, we can see the volume mounted at /tmp/data and are able to read the data created in the previous step:

```
$ kubectl exec -it sharevol -c c2 -- bash
[root@sharevol /]# mount | grep /tmp/data
/dev/vda3 on /tmp/data type xfs (rw,relatime,seclabel,attr2,inode64,prjquota)
[root@sharevol /]# cat /tmp/data/data/
cat: /tmp/data/data/: Not a directory
[root@sharevol /]# cat /tmp/data/data
some data
```

Note that in each container you need to decide where to mount the volume and that for emptyDir you currently can not specify resource consumption limits.

You can remove the pod with:

```
$ kubectl delete pod/sharevol
```

As already described, this will destroy the shared volume and all its contents.

12. Persistent Volumes

Kubernetes persistent volumes by example

A [persistent volume](#) (PV) is a cluster-wide resource that you can use to store data in a way that it persists beyond the lifetime of a pod. The PV is not backed by locally-attached storage on a worker node but by networked storage system such as *EBS* or *NFS* or a distributed filesystem like *Ceph*.

Create one persistent volume first using:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/pv/pv.yaml
```

In order to use a PV you need to claim it first, using a persistent volume claim (PVC). The PVC requests a PV with your desired specification (size, speed, etc.) from Kubernetes and binds it then to a pod where you can mount it as a volume. So let's create such a [PVC](#), asking Kubernetes for 1 GB of storage using the default [storage class](#):

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/pv/pvc.yaml

$ kubectl get pvc
NAME          STATUS    VOLUME                                     CAPACITY   ACCESS MODES
STORAGECLASS  AGE
myclaim       Bound    pvc-27fed6b6-3047-11e9-84bb-12b5519f9b58  1Gi        RW0
gp2-encrypted 18m
```

To understand how the persistency plays out, let's create a [deployment](#) that uses above PVC to mount it as a volume into `/tmp/persistent`:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/pv/deploy.yaml
```

Now we want to test if data in the volume actually persists. For this we find the pod managed by above deployment, exec into its main container and create a file called data in the `/tmp/persistent` directory (where we decided to mount the PV):

```
$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
pv-deploy-69959dccb5-jhxx         1/1     Running   0           16m

$ kubectl exec -it pv-deploy-69959dccb5-jhxxw -- bash
bash-4.2$ touch /tmp/persistent/data
bash-4.2$ ls /tmp/persistent/
data  lost+found
```

It's time to destroy the pod and let the deployment launch a new pod. The expectation is that the PV is available again in the new pod and the data in /tmp/persistent is still present. Let's check that:

```
$ kubectl delete po pv-deploy-69959dccb5-jhxxw
pod pv-deploy-69959dccb5-jhxxw deleted

$ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
pv-deploy-69959dccb5-kwrrv         1/1     Running   0           16m

$ kubectl exec -it pv-deploy-69959dccb5-kwrrv -- bash
bash-4.2$ ls /tmp/persistent/
data  lost+found
```

And indeed, the data file and its content is still where it is expected to be.

Note that the default behavior is that even when the deployment is deleted, the PVC (and the PV) continues to exist. This storage protection feature helps avoiding data loss. Once you're sure you don't need the data anymore, you can go ahead and delete the PVC and with it eventually destroy the PV:

```
$ kubectl delete pvc myclaim
persistentvolumeclaim "myclaim" deleted
```

The types of PV available in your Kubernetes cluster depend on the environment (on-premise or public cloud). Check out the [Stateful Kubernetes](https://stateful.kubernetes.sh/#storage) reference site if you want to learn more about this topic :

<https://stateful.kubernetes.sh/#storage>

13. StatefulSet

If you have a stateless app you want to use a deployment. However, for a stateful app you might want to use a [StatefulSet](#). Unlike a deployment, the StatefulSet provides certain guarantees about the identity of the pods it is managing (that is, predictable names) and about the startup order. Two more things that are different compared to a deployment: for network communication you need to create a [headless services](#) and for persistency the StatefulSet manages a [persistent volume](#) per pod.

In order to see how this all plays together, we will be using an [educational Kubernetes-native NoSQL datastore](#).

Let's start with creating the [stateful app](#), that is, the StatefulSet along with the persistent volumes and the headless service:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/mehdb/master/app.yaml
```

After a minute or so, you can have a look at all the resources that have been created:

```
$ kubectl get sts,po,pvc,svc
```

NAME	DESIRED	CURRENT	AGE
statefulset.apps/mehdb	2	2	1m

NAME	READY	STATUS	RESTARTS	AGE
pod/mehdb-0	1/1	Running	0	1m
pod/mehdb-1	1/1	Running	0	56s

NAME	STATUS	VOLUME
persistentvolumeclaim/data-mehdb-0	Bound	pvc-bc2d9b3b-310d-11e9-aeff-123713f594ec
1Gi		
persistentvolumeclaim/data-mehdb-1	Bound	pvc-d4b7620f-310d-11e9-aeff-123713f594ec
1Gi		

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/mehdb	ClusterIP	None	<none>	9876/TCP	1m

Now we can check if the stateful app is working properly. To do this, we use the `/status` endpoint of the headless service `mehdb:9876` and since we haven't put any data yet into the datastore, we'd expect that 0 keys are reported:

```
$ kubectl run -it --rm jumpod --restart=Never --image=quay.io/openshiftlabs/jump:0.2 -- curl mehdb:9876/status?level=full
```

If you don't see a command prompt, try pressing enter.

```
0  
pod "jumpod" deleted
```

And indeed we see 0 keys being available, reported above.

Note that sometimes a `StatefulSet` is not the best fit for your stateful app. You might be better off defining a [custom resource](#) along with writing a custom controller to have finer-grained control over your workload.

```
hljs.initHighlightingOnLoad(); renderMathInElement(document.body); var doNotTrack  
= false;  
if (!doNotTrack) {  
    (function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function(){  
        (i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new Date();a=s.createElement(o),  
        m=s.getElementsByTagName(o)  
    [0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)  
  
    })(window,document,'script','https://www.google-analytics.com/analytics.js','ga');  
    ga('create', 'UA-34425776-2', 'auto');  
  
    ga('send', 'pageview');  
}
```


14. Jobs

A **job** in Kubernetes is a supervisor for pods carrying out batch processes, that is, a process that runs for a certain time to completion, for example a calculation or a backup operation.

Let's create a **job** called **countdown** that supervises a pod counting from 9 down to 1:

```
$ kubectl apply -f https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/jobs/job.yaml
```

You can see the job and the pod it looks after like so:

```
$ kubectl get jobs
NAME                DESIRED    SUCCESSFUL    AGE
countdown           1          1             5s

$ kubectl get pods
NAME                READY    STATUS      RESTARTS    AGE
countdown-qkjsx8    0/1     Completed   0           2m17s
```

To learn more about the status of the job, do:

```
$ kubectl describe jobs/countdown
Name:                countdown
Namespace:           default
Selector:             controller-uid=4960c8be-6a3f-11ea-84fd-0242ac11002a
Labels:              controller-uid=4960c8be-6a3f-11ea-84fd-0242ac11002a
                    job-name=countdown
Annotations:         kubectl.kubernetes.io/last-applied-configuration:
                    {"apiVersion":"batch/v1","kind":"Job","metadata":{"annotations":
{"},"name":"countdown","namespace":"default"},"spec":{"template":{"metadata...
Parallelism:         1
Completions:         1
Start Time:          Fri, 20 Mar 2020 00:11:03 +0000
Completed At:        Fri, 20 Mar 2020 00:11:12 +0000
Duration:            9s
Pods Statuses:       0 Running / 1 Succeeded / 0 Failed
Pod Template:
  Labels:  controller-uid=4960c8be-6a3f-11ea-84fd-0242ac11002a
          job-name=countdown
  Containers:
    counter:
      Image:      centos:7
      Port:       <none>
      Host Port:  <none>
```

```

Command:
  bin/bash
  -c
  for i in 9 8 7 6 5 4 3 2 1 ; do echo $i ; done
Environment:  <none>
Mounts:       <none>
Volumes:      <none>
Events:
  Type           Reason             Age           From           Message
  ----           -
Normal          SuccessfulCreate    3m18s         job-controller  Created pod: countdown-qkjsx8

```

And to see the output of the job via the pod it supervised, execute:

```
kubectl logs countdown-qkjsx8
```

```

9
8
7
6
5
4
3
2
1

```

To clean up, use the delete verb on the job object which will remove all the supervised pods:

```

$ kubectl delete job countdown
job "countdown" deleted

```

Note that there are also more advanced ways to use jobs, for example, by utilizing a [work queue](#) or scheduling the execution at a certain time via [cron jobs](#).

```

hljs.initHighlightingOnLoad(); renderMathInElement(document.body); var
doNotTrack = false;
if (!doNotTrack) {
  (function(i,s,o,g,r,a,m){i['GoogleAnalyticsObject']=r;i[r]=i[r]||function()
{
  (i[r].q=i[r].q||[]).push(arguments)},i[r].l=1*new
Date();a=s.createElement(o),
  m=s.getElementsByTagName(o)
[0];a.async=1;a.src=g;m.parentNode.insertBefore(a,m)

})(window,document,'script','https://www.google-analytics.com/analytics.js','
ga');

```

```
ga('create', 'UA-34425776-2', 'auto');  
  
ga('send', 'pageview');  
}
```